

Self-Play Scaling Experiments: Per-Iteration Results

The full empirical record behind the draw-cycle study — every training iteration, evaluation match, and Elo estimate from the 9.8-hour GPU run.

Daniel Lichtenberger · July 2026
[rlchess.xyz](#) · [Source on GitHub](#)

Hardware: Kaggle, 2× Tesla T4 (GPU) · **Driver:** `python run_experiment.py --device cuda` **Date:** 2026-07-14 · **Training wall-clock:** ~9.8 h (baseline 68.9 min + scaled 444.7 min + assisted 75.2 min)

Provenance. The Kaggle notebook's automatic `git push` failed with a 403 (its `GITHUB_TOKEN` secret lacked write scope), so every number in this file is transcribed directly from that run's captured stdout — the per-iteration self-play termination tables, the periodic evaluation lines, and the loss curve. Two things could **not** be recovered from the surviving logs and are marked where they appear: (1) the win/draw/loss split of each stage's *final* stand-alone 50-game evaluation (only its Elo survived), and (2) a downloadable trained checkpoint, which was lost with the Kaggle session (re-run the driver to regenerate one). The figures below are plotted directly from the per-iteration numbers transcribed in this file.

 The draw cycle — non-decisive rate and Elo vs. random across the three stages

Run configuration

All three stages share the identical code and network (4 residual blocks, 64 channels); they differ only in the variable under test.

Stage	Iterations	Games/iter	Total games	Sims/move	Self-play material w
baseline	10	10	100	100	0.00
scaled	25	40	1,000	160	0.00
assisted	10	10	100	100	0.30

Evaluation at each checkpoint is 50 games vs. `RandomAgent` at 50 sims/move.

Headline results

The driver's two headline metrics are the self-play **non-decisive rate** (fraction of games *not* ending in checkmate) and **Elo vs. random**. "Non-decisive rate" below is the **final iteration's** rate (the end-of-

training state); the aggregate over *all* of a stage's games is given alongside because for the baseline the two differ sharply (see §A).

Experiment A — compute scaling (material $w = 0$ in both)

Run	Total games	Sims/move	Non-decisive (final iter / all games)	Elo vs random
baseline	100	100	100% / 92%	-21
scaled	1,000	160	92% / 88%	-7

Result. 10× the self-play raises Elo ($-21 \rightarrow -7$) and produces more checkmates in aggregate ($8/100 = 8\% \rightarrow 120/1000 = 12\%$), but the cycle is not broken: the network still cannot reliably convert, and Elo stays negative.

Experiment B — self-play signal at fixed compute (same budget as baseline)

Run	Total games	Sims/move	Material w	Non-decisive (final iter / all games)	Elo vs random
baseline	100	100	0.00	100% / 92%	-21
assisted	100	100	0.30	70% / 81%	+28

Result. Blending a 0.30 material term into self-play leaf evaluation — at the *same* compute as the baseline — drops the final-iteration non-decisive rate to 70%, raises the aggregate checkmate count to 19/100, and flips Elo from -21 to +28. The engine now beats random. This is the more cost-effective lever.

Reading the two together

Both levers move the metrics, so the plateau is driven by **both** limited compute *and* an impoverished self-play signal (the "yes | yes" row of the §5.3 reading matrix) — but the signal lever does more per unit of compute.

A. Per-iteration self-play termination

Counts are games ending by each reason. **CM** checkmate · **ST** stalemate · **INSUF** insufficient material · **REP** threefold repetition · **50M** fifty-move · **CAP** hit the `max_moves = 200` cap. "Non-dec %" = fraction not won by checkmate.

Baseline — 10 iters × 10 games, 100 sims, mat-w 0.00

Iter	CM	ST	INSUF	REP	50M	CAP	Non-dec %
1	6	0	0	0	0	4	40%
2	2	0	0	8	0	0	80%
3	0	0	0	10	0	0	100%
4	0	0	0	10	0	0	100%
5	0	0	0	10	0	0	100%
6	0	0	0	10	0	0	100%
7	0	0	0	10	0	0	100%
8	0	0	0	10	0	0	100%
9	0	0	0	10	0	0	100%
10	0	0	0	10	0	0	100%
ALL	8	0	0	88	0	4	92%

*The draw cycle forming in real time: the early, near-random network scores a few accidental checkmates (8 in iters 1–2), then collapses into pure threefold repetition from iter 3 onward — zero checkmates for the last 80 games. Draws are driven by **repetition**, not the fifty-move rule.*

Scaled — 25 iters × 40 games, 160 sims, mat-w 0.00

Iter	CM	ST	INSUF	REP	50M	CAP	Non-dec %
1	29	0	0	0	0	11	27.5%
2	11	0	0	29	0	0	72.5%
3	10	0	0	30	0	0	75.0%
4	5	0	0	34	0	1	87.5%
5	9	0	0	31	0	0	77.5%
6	4	0	0	36	0	0	90.0%
7	4	0	0	36	0	0	90.0%
8	3	0	0	37	0	0	92.5%
9	4	0	0	35	0	1	90.0%
10	2	0	0	38	0	0	95.0%
11	5	0	0	35	0	0	87.5%
12	4	0	0	36	0	0	90.0%
13	4	0	0	36	0	0	90.0%
14	0	0	0	39	0	1	100.0%
15	2	0	0	38	0	0	95.0%
16	1	0	0	39	0	0	97.5%
17	4	0	0	36	0	0	90.0%
18	4	0	0	36	0	0	90.0%
19	1	0	0	39	0	0	97.5%
20	3	0	0	37	0	0	92.5%
21	2	0	0	38	0	0	95.0%
22	5	0	0	35	0	0	87.5%
23	1	0	0	39	0	0	97.5%
24	0	0	0	40	0	0	100.0%
25	3	0	0	37	0	0	92.5%
ALL	120	0	0	866	0	14	88%

Same pattern at scale: iter 1's untrained net delivers 29 accidental mates, then the rate settles into the high-80s/90s% non-decisive band for the rest of training. More compute keeps a residual trickle of checkmates alive (120/1000 aggregate vs. the baseline's 8/100) but does not escape the plateau.

Assisted — 10 iters × 10 games, 100 sims, mat-w 0.30

Iter	CM	ST	INSUF	REP	50M	CAP	Non-dec %
1	5	0	2	1	0	2	50%
2	2	0	0	7	0	1	80%
3	1	0	0	9	0	0	90%
4	1	0	0	9	0	0	90%
5	2	0	0	8	0	0	80%
6	0	0	0	10	0	0	100%
7	2	0	0	7	0	1	80%
8	2	0	0	7	0	1	80%
9	1	0	0	9	0	0	90%
10	3	0	0	7	0	0	70%
ALL	19	0	2	74	0	5	81%

With the material assist, checkmates keep appearing throughout training instead of dying out — 19/100 aggregate, and the final iteration hits 30% checkmate (70% non-decisive), the best of any stage. Decisive games give the value head a real signal to learn from.

B. Elo vs. random — periodic evaluation checkpoints

Each row is an independent 50-game match vs. `RandomAgent` (50 sims/move), logged during training at the stage's `eval_every` interval. W/D/L are the raw counts.

Baseline (every 2 iters)

Iter	W	D	L	Score	Elo
2	2	43	5	47.0%	-21
4	0	44	6	44.0%	-42
6	3	42	5	48.0%	-14
8	1	45	4	47.0%	-21
10	2	44	4	48.0%	-14

Scaled (every 5 iters)

Iter	W	D	L	Score	Elo
5	8	40	2	56.0%	+42
10	8	39	3	55.0%	+35
15	0	47	3	47.0%	-21
20	6	40	4	52.0%	+14
25	3	41	6	47.0%	-21

Assisted (every 2 iters)

Iter	W	D	L	Score	Elo
2	6	39	5	51.0%	+7
4	6	40	4	52.0%	+14
6	5	42	3	52.0%	+14
8	7	40	3	54.0%	+28
10	10	36	4	56.0%	+42

Final stand-alone evaluation (a separate 50-game match after training; only the Elo survives in the logs — the W/D/L split was not captured): baseline **-21**, scaled **-7**, assisted **+28**. These are the headline Elo figures above. Note they differ from each stage's *last periodic* checkpoint (e.g. scaled's iter-25 checkpoint read -21 but its final match read -7) — independent 50-game samples of a near-even engine are noisy, which is itself worth reporting.

C. Training loss

 Training dynamics of the scaled run — policy, value, and total loss

Total loss = policy cross-entropy + value MSE. Start (iter 1) → end (final iter):

Stage	Policy loss (start → end)	Value loss (start → end)	Total (start → end)
baseline	4.54 → 0.81	0.092 → 0.006	4.63 → 0.81
scaled	4.02 → 0.96	0.164 → 0.012	4.19 → 0.97
assisted	4.37 → 0.69	0.101 → 0.023	4.47 → 0.71

Policy loss falls sharply in every stage — the network readily learns to imitate its own MCTS visit distribution. The value loss is tiny from the start and shrinks toward ≈ 0.01 : with almost every game drawn, the

outcome target `z ≈ 0` nearly everywhere, so the value head has almost nothing to fit. That collapse is the draw-cycle signature (technical report §6), not a training success.

Key findings

1. **The draw cycle is real and self-reinforcing.** At baseline compute the near-random network scores a few accidental mates in iters 1–2, then collapses to 100% repetition draws from iter 3 on — the value target flattens to ≈ 0 and the cycle locks in.
 2. **Compute helps, but does not break it.** 10× the self-play (100 → 1,000 games) lifts Elo (−21 → −7) and keeps a residual checkmate rate alive (8% → 12% aggregate), yet Elo stays negative and the non-decisive rate stays in the high 80s/90s%.
 3. **Signal quality is the more efficient lever.** `material_weight = 0.30` in self-play — at the *same* baseline compute — drops the final-iteration non-decisive rate to 70%, keeps checkmates appearing throughout training (19/100 aggregate), and flips Elo to +28. The engine beats random.
 4. **Both effects compound** ("yes | yes" in the §5.3 reading matrix): compute helps *and* signal helps. A combined scaled-plus-assisted run is the natural next step.
 5. `material_weight = 0.30` is the recommended self-play default, now set via `MCTConfig.material_weight` and exposed as `run_experiment.py --assist-weight`.
-

Reproducibility

```
# All three stages on a GPU (this run):
python run_experiment.py --device cuda

# Fast pipeline check on CPU (~1-2 min, results not meaningful):
python run_experiment.py --quick
```

Config used (from `config.py` ; stages differ only in iterations/games/sims and the self-play material weight noted above):

- Network: `num_residual_blocks = 4`, `num_channels = 64`
 - Search: `c_puct = 1.5`, `dirichlet_alpha = 0.3`, `dirichlet_epsilon = 0.25`
 - Self-play: `temperature_moves = 15`, `max_moves = 200`
 - Training: `epochs_per_iteration = 4`, `batch_size = 64`, `learning_rate = 1e-3`, `weight_decay = 1e-4`, `replay_buffer_size = 20 000`
-

Every figure in this report is reproduced from the project's own test suite, experiment logs, or committed results files. Where a result failed to beat its baseline, that is what is reported. · Source and full history: github.com/Danny-397